


URPP tutorial on dplyr/tidyverse

Stefan Wyder - URPP Evolution in Action

June 1, 2018



R for Data Science



<http://r4ds.had.co.nz>

Installation

You can install the complete tidyverse with a single line of code:

```
install.packages("tidyverse")
```

install also the data package nycflights13

```
install.packages("nycflights13")
```

Load packages

```
library(tidyverse)  
library(nycflights13)
```

This data frame contains all 336,776 flights that departed from New York City in 2013 flights
see ?flights

Data Science cycle

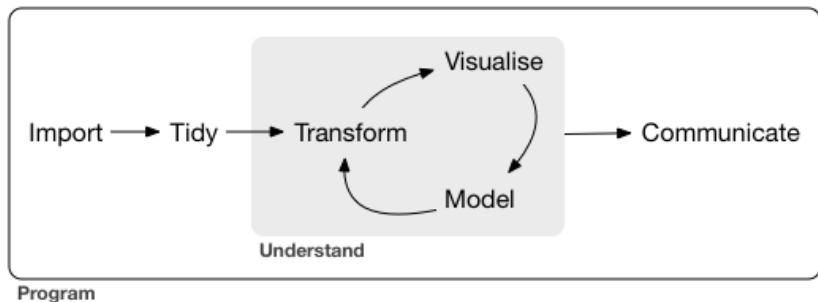


Table of Contents

- Data transformation (Chapter 5)
- Tibbles (Chapter 10)
- Relational data (Chapter 13)

Tidy data

There are three interrelated rules which make a dataset tidy:

- 1 Each variable must have its own column.
- 2 Each observation must have its own row.
- 3 Each value must have its own cell.

country	year	cases	population
Afghanistan	1999	1815	19787071
Afghanistan	2000	2666	20595360
Brazil	1999	31737	172006362
Brazil	2000	80488	174004898
China	1999	212258	1272015272
China	2000	213766	1280008583

variables

country	year	cases	population
Afghanistan	1999	1815	19787071
Afghanistan	2000	2666	20595360
Brazil	1999	31737	172006362
Brazil	2000	80488	174004898
China	1999	212258	1272015272
China	2000	213766	1280008583

observations

country	year	cases	population
Afghanistan	1999	1815	19787071
Afghanistan	2000	2666	20595360
Brazil	1999	31737	172006362
Brazil	2000	80488	174004898
China	1999	212258	1272015272
China	2000	213766	1280008583

values

Tibble (Chapter 10)

Tibbles are data frames, but slightly tweaked to work better in the tidyverse

- better printing
- never changes the type of the inputs (strings to factors!)
- never changes the names of variables

```
flights
```

```
## # A tibble: 336,776 x 19
##   year month   day dep_time sched_dep_time dep_delay arr_time
##   <int> <int> <int>   <int>         <int>         <dbl>   <int>
## 1  2013     1     1     517             515           2.     830
## 2  2013     1     1     533             529           4.     850
## 3  2013     1     1     542             540           2.     923
## 4  2013     1     1     544             545          -1.    1004
## 5  2013     1     1     554             600          -6.     812
## 6  2013     1     1     554             558          -4.     740
## 7  2013     1     1     555             600          -5.     913
## 8  2013     1     1     557             600          -3.     709
## 9  2013     1     1     557             600          -3.     838
## 10 2013     1     1     558             600          -2.     753
## # ... with 336,766 more rows, and 12 more variables: sched_arr_time <int>,
## #   arr_delay <dbl>, carrier <chr>, flight <int>, tailnum <chr>,
## #   origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>, hour <dbl>,
## #   minute <dbl>, time_hour <dtm>
```

5 key dplyr functions

- `filter()`: Pick observations by their values
- `arrange()`: Reorder the rows
- `select()`: Pick variables by their names
- `mutate()`: Create new variables with functions of existing variables
- `summarise()`: Collapse many values down to a single summary

1 example

```
filter(flights, month == 1, day == 1)
```

```
## # A tibble: 842 x 19
##   year month   day dep_time sched_dep_time dep_delay arr_time
##   <int> <int> <int>   <int>         <int>         <dbl>   <int>
## 1  2013     1     1     517             515           2.     830
## 2  2013     1     1     533             529           4.     850
## 3  2013     1     1     542             540           2.     923
## 4  2013     1     1     544             545          -1.    1004
## 5  2013     1     1     554             600          -6.     812
## 6  2013     1     1     554             558          -4.     740
## 7  2013     1     1     555             600          -5.     913
## 8  2013     1     1     557             600          -3.     709
## 9  2013     1     1     557             600          -3.     838
## 10 2013     1     1     558             600          -2.     753
## # ... with 832 more rows, and 12 more variables: sched_arr_time <int>,
## #   arr_delay <dbl>, carrier <chr>, flight <int>, tailnum <chr>,
## #   origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>, hour <dbl>,
## #   minute <dbl>, time_hour <dtm>
```

filter()








Guess which ones are correct. What does it do?

```
filter(flights, month == 1, day == 1)
filter(flights, month == 11 & month == 12)
filter(flights, month == 11 | month == 12)
filter(flights, month == 11 | 12)
filter(flights, month %in% c(11, 12))
```

filter()

```
# finds all flights that departed on Jan 1
filter(flights, month == 1, day == 1)
filter(flights, month == 1 & day == 1)
# finds all flights that departed in Nov or Dec:
filter(flights, month == 11 | month == 12)
# doesn't work:
filter(flights, month == 11 | 12)
# %in% also works:
filter(flights, month %in% c(11, 12))
```

Boolean operations



Be careful with arithmetic comparisons

```
sqrt(2) ^ 2 == 2
```

```
## [1] FALSE
```

```
1/49 * 49 == 1
```

```
## [1] FALSE
```

Use `near` with a tolerance:

```
near(sqrt(2) ^ 2, 2)
```

```
## [1] TRUE
```

```
near(1 / 49 * 49, 1)
```

```
## [1] TRUE
```

Exercise 1

Find flights that weren't delayed (on arrival or departure) by more than two hours

Solution 1

You could use either of the following two filters:

```
filter(flights, arr_delay <= 120, dep_delay <= 120)  
filter(flights, !(arr_delay > 120 | dep_delay > 120))
```

Missing values

`filter()` only includes rows where the condition is TRUE; it excludes both FALSE and NA values.

If you want to preserve missing values, ask for them explicitly:

```
filter(df, is.na(x) | x > 1)
```

Exercises

5.2.4 Exercises

arrange(): order rows

arrange() works similarly to filter() except that instead of selecting rows, it changes their order.

```
arrange(flights, year)
arrange(flights, year, month, day)
arrange(flights, desc(arr_delay)) # order in decreasing order
```

missing values are always sorted at the end

Exercises

5.3.1 Exercises

Select columns with select()

```
select(flights, year, month, day)
```

```
# Select all columns between year and day (inclusive)
```

```
select(flights, year:day)
```

everything() is useful if you have a handful of variables you'd like to move to the start of the data frame

```
select(flights, time_hour, air_time, everything())
```

select()

There are a number of helper functions you can use within `select()`:

- `starts_with("abc")`: matches names that begin with "abc".
- `ends_with("xyz")`: matches names that end with "xyz".
- `contains("ijk")`: matches names that contain "ijk".
- `matches("(.)\\1")`: selects variables that match a regular expression. Here: any variables that contain repeated characters.
- `num_range("x", 1:3)` matches `x1`, `x2` and `x3`.

See `?select` for more details.

Exercises

5.4.1 Exercises

Add new variables with mutate()

mutate() always adds new columns at the end of your dataset, as function of the existing columns.

```
flights_sml <- select(flights,
  year:day,
  ends_with("delay"),
  distance,
  air_time
)
mutate(flights_sml,
  gain = arr_delay - dep_delay,
  hours = air_time / 60,
  gain_per_hour = gain / hours
)
```

Note that you can refer to columns that you've just created

`mutate()` works with all vectorized functions

The key property is that the function must be vectorised: it must take a vector of values as input, return a vector with the same number of values as output.

- Arithmetic operators: `+` `-` `*` `/` `^`
- Modular arithmetic: `%/%` (integer division) and `%%` (remainder)
- Logs: `log()`, `log2()`, `log10()`
- Offsets: `lead()` and `lag()`
- Logical comparisons: `<`, `<=`, `>`, `>=`, `!=`
- Ranking: `min_rank()`, `min_rank(desc())` also its variants `row_number()`, `dense_rank()`, `percent_rank()`, `cume_dist()`, `ntile()`
- Cumulative and rolling aggregates: running sums, products, mins, maxes and means: `cumsum()`, `cumprod()`, `cummin()`, `cummax()` and `cummean()`
- conjunction of functions: i.e. `x / sum(x)`: proportion of a total, and `y - mean(y)`: diff from the mean

transmute()

If you only want to keep the new variables, use `transmute()`:

```
transmute(flights,  
  dep_time,  
  hour = dep_time %/% 100,  
  minute = dep_time %% 100  
)
```

`%/%` (integer division) handy to because it allows you to break integers up into pieces (here: hours)

Exercises

5.5.2 Exercises

Grouped summaries with summarise()

summarise() is not terribly useful unless we pair it with group_by():

```
by_day <- group_by(flights, year, month, day)
summarise(by_day, delay = mean(dep_delay, na.rm = TRUE))
```

```
## # A tibble: 365 x 4
## # Groups:   year, month [?]
##   year month   day delay
##   <int> <int> <int> <dbl>
## 1  2013     1     1  11.5
## 2  2013     1     2  13.9
## 3  2013     1     3  11.0
## 4  2013     1     4   8.95
## 5  2013     1     5   5.73
## 6  2013     1     6   7.15
## 7  2013     1     7   5.42
## 8  2013     1     8   2.55
## 9  2013     1     9   2.28
## 10 2013     1    10   2.84
## # ... with 355 more rows
```

Combining multiple operations with the pipe

Explore the relationship between the distance and average delay for each location

```
by_dest <- group_by(flights, dest)
delay <- summarise(by_dest,
  count = n(),
  dist = mean(distance, na.rm = TRUE),
  delay = mean(arr_delay, na.rm = TRUE)
)
delay <- filter(delay, count > 20, dest != "HNL")
```

a little frustrating to write because we have to give each intermediate data frame a name, even though we don't care about it. Naming things is hard and slow.

the pipe %>%

There's another way to tackle the same problem with the pipe, %>%:

```
delays <- flights %>%  
  group_by(dest) %>%  
  summarise(  
    count = n(),  
    dist = mean(distance, na.rm = TRUE),  
    delay = mean(arr_delay, na.rm = TRUE)  
  ) %>%  
  filter(count > 20, dest != "HNL")
```

- pronounced as “then”
- much more readable (left-right, top-bottom)
- key criteria for belonging to the tidyverse (only exception: ggplot2)

Missing values

```
flights %>%  
  group_by(year, month, day) %>%  
  summarise(mean = mean(dep_delay))
```

```
## # A tibble: 365 x 4  
## # Groups:   year, month [?]  
##   year month   day mean  
##   <int> <int> <int> <dbl>  
## 1  2013     1     1    NA  
## 2  2013     1     2    NA  
## 3  2013     1     3    NA  
## 4  2013     1     4    NA  
## 5  2013     1     5    NA  
## 6  2013     1     6    NA  
## 7  2013     1     7    NA  
## 8  2013     1     8    NA  
## 9  2013     1     9    NA  
## 10 2013     1    10    NA  
## # ... with 355 more rows
```

We get a lot of missing values! Usual rule of missing values: if there's any missing value in the input, the output will be a missing value.

Missing values cont.

Fortunately, all aggregation functions have an `na.rm` argument which removes the missing values prior to computation:

```
flights %>%  
  group_by(year, month, day) %>%  
  summarise(mean = mean(dep_delay, na.rm = TRUE))
```

```
## # A tibble: 365 x 4  
## # Groups:   year, month [?]  
##   year month   day mean  
##   <int> <int> <int> <dbl>  
## 1  2013     1     1 11.5  
## 2  2013     1     2 13.9  
## 3  2013     1     3 11.0  
## 4  2013     1     4  8.95  
## 5  2013     1     5  5.73  
## 6  2013     1     6  7.15  
## 7  2013     1     7  5.42  
## 8  2013     1     8  2.55  
## 9  2013     1     9  2.28  
## 10 2013     1    10  2.84  
## # ... with 355 more rows
```

Missing values are due to cancelled flights. In this case we could also create a new data frame without cancelled flights.

Counts

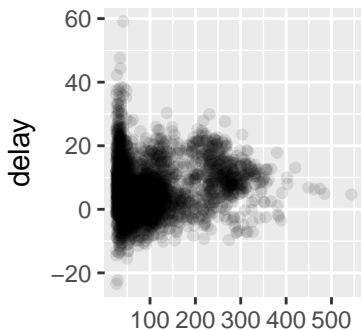
Whenever you do any aggregation, it's always a good idea to include either a count `n()`, or a count of non-missing values `sum(!is.na(x))`. Not to draw conclusions based on very small amounts of data

```
not_cancelled <- flights %>%  
  filter(!is.na(dep_delay), !is.na(arr_delay))  
delays <- not_cancelled %>%  
  group_by(tailnum) %>%  
  summarise(  
    delay = mean(arr_delay, na.rm = TRUE),  
    n = n()  
  )
```

Plot

It's often useful to filter out the groups with the smallest numbers of observations, so you can see more of the pattern and less of the extreme variation in the smallest groups

```
delays %>%  
  filter(n > 25) %>%  
  ggplot(mapping = aes(x = n, y = delay)) +  
    geom_point(alpha = 1/10)
```



Variation

- Measures of location: `mean(x)`, `median(x)`
- Measures of spread: `sd(x)`, `IQR(x)`, `mad(x)`
- Measures of rank: `min(x)`, `quantile(x, 0.25)`, `max(x)`
- Measures of position: `first(x)`, `nth(x, 2)`, `last(x)`: allow to set a default position
- Counts: `n()` returns the size of the current group.
`sum(!is.na(x))` returns the number of non-missing values
`n_distinct(x)` counts the number of distinct (unique) values

Counts

Counts are so useful that dplyr provides a simple helper if all you want is a count:

```
not_cancelled %>%  
  count(dest)
```

```
## # A tibble: 104 x 2  
##   dest      n  
##   <chr> <int>  
## 1 ABQ     254  
## 2 ACK     264  
## 3 ALB     418  
## 4 ANC       8  
## 5 ATL    16837  
## 6 AUS     2411  
## 7 AVL     261  
## 8 BDL     412  
## 9 BGR     358  
## 10 BHM     269  
## # ... with 94 more rows
```

Counts cont.

You can optionally provide a weight variable. For example, you could use this to “count” (sum) the total number of miles a plane flew:

```
not_cancelled %>%  
  count(tailnum, wt = distance)
```

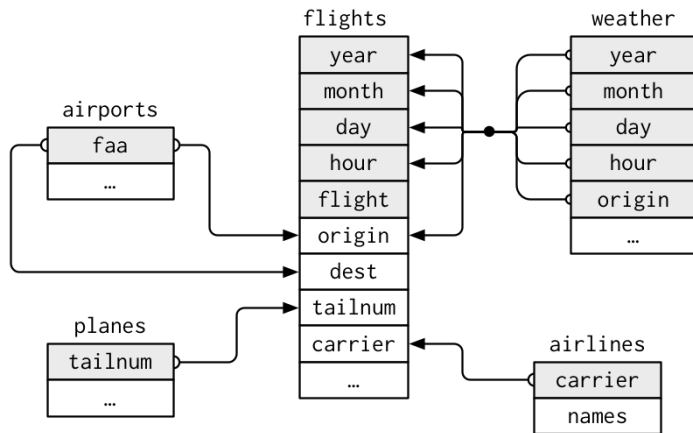
```
## # A tibble: 4,037 x 2  
##   tailnum      n  
##   <chr>    <dbl>  
## 1 D942DN    3418.  
## 2 NOEGMQ 239143.  
## 3 N10156 109664.  
## 4 N102UW  25722.  
## 5 N103US  24619.  
## 6 N104UW  24616.  
## 7 N10575 139903.  
## 8 N105UW  23618.  
## 9 N107US  21677.  
## 10 N108UW  32070.  
## # ... with 4,027 more rows
```

Exercises

5.6.7 Exercises

Relational data (Chapter 13)

Collectively, multiple tables of data are called **relational data** (relations between a pair of tables)



Relational data cont.

- Mutating joins, which add new variables to one data frame from matching observations in another.
- Filtering joins, which filter observations from one data frame based on whether or not they match an observation in the other table.
- Set operations, which treat observations as if they were set elements.

Generally, dplyr is a little easier to use than SQL because dplyr is specialised to do data analysis.

Other packages of the tidyverse

- readr: import data (flat files)

Other packages of the tidyverse

- readr: import data (flat files)
- tidyr: tidy data

Other packages of the tidyverse

- readr: import data (flat files)
- tidyr: tidy data
- dplyr: wrangle data

Other packages of the tidyverse

- readr: import data (flat files)
- tidyr: tidy data
- dplyr: wrangle data
- stringr: functions to work with strings

Other packages of the tidyverse

- readr: import data (flat files)
- tidyr: tidy data
- dplyr: wrangle data
- stringr: functions to work with strings
- forcats: functions to work with factors

Other packages of the tidyverse

- readr: import data (flat files)
- tidyr: tidy data
- dplyr: wrangle data
- stringr: functions to work with strings
- forcats: functions to work with factors
- purrr: iteration, replacement of apply functions

Other packages of the tidyverse

- readr: import data (flat files)
- tidyr: tidy data
- dplyr: wrangle data
- stringr: functions to work with strings
- forcats: functions to work with factors
- purrr: iteration, replacement of apply functions
- ggplot2: plot data

Other packages of the tidyverse

- readr: import data (flat files)
 - tidyr: tidy data
 - dplyr: wrangle data
 - stringr: functions to work with strings
 - forcats: functions to work with factors
 - purrr: iteration, replacement of apply functions
 - ggplot2: plot data
-
- even more packages: e.g. import .xls and .xlsx sheets (readxl),
for dates and date-times (lubridate), ...