

Solutions for dplyr/tidyverse URPP tutorial.Rmd

Stefan Wyder

June 4, 2018

```
library(tidyverse)
library(nycflights13)
```

Exercises 2

Find all flights that

1. Had an arrival delay of two or more hours

```
filter(flights, arr_delay >= 120)
```

```
## # A tibble: 10,200 x 19
##   year month   day dep_time sched_dep_time dep_delay arr_time
##   <int> <int> <int>   <int>         <int>      <dbl>   <int>
## 1  2013     1     1     811           630        101.    1047
## 2  2013     1     1     848           1835       853.    1001
## 3  2013     1     1     957           733        144.    1056
## 4  2013     1     1    1114           900        134.    1447
## 5  2013     1     1    1505          1310       115.    1638
## 6  2013     1     1    1525          1340       105.    1831
## 7  2013     1     1    1549          1445        64.    1912
## 8  2013     1     1    1558          1359       119.    1718
## 9  2013     1     1    1732          1630        62.    2028
## 10 2013     1     1    1803          1620       103.    2008
## # ... with 10,190 more rows, and 12 more variables: sched_arr_time <int>,
## #   arr_delay <dbl>, carrier <chr>, flight <int>, tailnum <chr>,
## #   origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>, hour <dbl>,
## #   minute <dbl>, time_hour <dtm>
```

2. Flew to Houston (IAH or HOU)

```
filter(flights, dest %in% c("IAH", "HOU"))
```

```
## # A tibble: 9,313 x 19
##   year month   day dep_time sched_dep_time dep_delay arr_time
##   <int> <int> <int>   <int>         <int>      <dbl>   <int>
## 1  2013     1     1     517           515         2.     830
## 2  2013     1     1     533           529         4.     850
## 3  2013     1     1     623           627        -4.     933
## 4  2013     1     1     728           732        -4.    1041
## 5  2013     1     1     739           739         0.    1104
## 6  2013     1     1     908           908         0.    1228
## 7  2013     1     1    1028          1026         2.    1350
## 8  2013     1     1    1044          1045        -1.    1352
## 9  2013     1     1    1114           900       134.    1447
## 10 2013     1     1    1205          1200         5.    1503
## # ... with 9,303 more rows, and 12 more variables: sched_arr_time <int>,
## #   arr_delay <dbl>, carrier <chr>, flight <int>, tailnum <chr>,
```

```
## #   origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>, hour <dbl>,
## #   minute <dbl>, time_hour <dtm>
```

3. Departed in summer (July, August, and September)

```
filter(flights, month >= 7, month <= 9)
```

```
## # A tibble: 86,326 x 19
##   year month   day dep_time sched_dep_time dep_delay arr_time
##   <int> <int> <int>   <int>         <int>         <dbl>   <int>
## 1  2013     7     1       1           2029         212.    236
## 2  2013     7     1       2           2359          3.    344
## 3  2013     7     1      29           2245        104.    151
## 4  2013     7     1     43           2130        193.    322
## 5  2013     7     1     44           2150        174.    300
## 6  2013     7     1     46           2051        235.    304
## 7  2013     7     1     48           2001        287.    308
## 8  2013     7     1     58           2155        183.    335
## 9  2013     7     1    100           2146        194.    327
##10  2013     7     1    100           2245        135.    337
## # ... with 86,316 more rows, and 12 more variables: sched_arr_time <int>,
## #   arr_delay <dbl>, carrier <chr>, flight <int>, tailnum <chr>,
## #   origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>, hour <dbl>,
## #   minute <dbl>, time_hour <dtm>
```

```
# also correct:
# filter(flights, between(month, 7, 9))
# filter(flights, month in c(7,8,9))
```

4. Arrived more than two hours late, but didn't leave late

```
filter(flights, dep_delay <= 0, arr_delay > 120)
```

```
## # A tibble: 29 x 19
##   year month   day dep_time sched_dep_time dep_delay arr_time
##   <int> <int> <int>   <int>         <int>         <dbl>   <int>
## 1  2013     1    27    1419           1420         -1.    1754
## 2  2013    10     7    1350           1350          0.    1736
## 3  2013    10     7    1357           1359         -2.    1858
## 4  2013    10    16     657           700         -3.    1258
## 5  2013    11     1     658           700         -2.    1329
## 6  2013     3    18    1844           1847         -3.     39
## 7  2013     4    17    1635           1640         -5.    2049
## 8  2013     4    18     558           600         -2.    1149
## 9  2013     4    18     655           700         -5.    1213
##10  2013     5    22    1827           1830         -3.    2217
## # ... with 19 more rows, and 12 more variables: sched_arr_time <int>,
## #   arr_delay <dbl>, carrier <chr>, flight <int>, tailnum <chr>,
## #   origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>, hour <dbl>,
## #   minute <dbl>, time_hour <dtm>
```

```
# these are also correct (the !is.na condition is more explicit but redundant):
# filter(flights, !is.na(dep_delay), dep_delay <= 0, arr_delay > 120)
# filter(flights, !is.na(dep_delay) & dep_delay <= 0 & arr_delay > 120)
```

5. Were delayed by at least an hour, but made up over 30 minutes in flight

```
filter(flights, dep_delay >= 60, dep_delay - arr_delay > 30)
```

```
## # A tibble: 1,844 x 19
##   year month   day dep_time sched_dep_time dep_delay arr_time
##   <int> <int> <int>   <int>         <int>         <dbl>   <int>
## 1  2013     1     1    2205           1720         285.     46
## 2  2013     1     1    2326           2130         116.    131
## 3  2013     1     3    1503           1221         162.   1803
## 4  2013     1     3    1839           1700          99.   2056
## 5  2013     1     3    1850           1745          65.   2148
## 6  2013     1     3    1941           1759         102.   2246
## 7  2013     1     3    1950           1845          65.   2228
## 8  2013     1     3    2015           1915          60.   2135
## 9  2013     1     3    2257           2000         177.     45
## 10 2013     1     4    1917           1700         137.   2135
## # ... with 1,834 more rows, and 12 more variables: sched_arr_time <int>,
## #   arr_delay <dbl>, carrier <chr>, flight <int>, tailnum <chr>,
## #   origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>, hour <dbl>,
## #   minute <dbl>, time_hour <dtm>

# explicit omission of NAs is redundant but might increase readability
# filter(flights, !is.na(dep_delay),
#         dep_delay >= 60, dep_delay - arr_delay > 30)
```

6. How many flights have a missing dep_time? What other variables are missing? What might these rows represent?

```
filter(flights, is.na(dep_time))
```

```
## # A tibble: 8,255 x 19
##   year month   day dep_time sched_dep_time dep_delay arr_time
##   <int> <int> <int>   <int>         <int>         <dbl>   <int>
## 1  2013     1     1     NA           1630          NA     NA
## 2  2013     1     1     NA           1935          NA     NA
## 3  2013     1     1     NA           1500          NA     NA
## 4  2013     1     1     NA            600          NA     NA
## 5  2013     1     2     NA           1540          NA     NA
## 6  2013     1     2     NA           1620          NA     NA
## 7  2013     1     2     NA           1355          NA     NA
## 8  2013     1     2     NA           1420          NA     NA
## 9  2013     1     2     NA           1321          NA     NA
## 10 2013     1     2     NA           1545          NA     NA
## # ... with 8,245 more rows, and 12 more variables: sched_arr_time <int>,
## #   arr_delay <dbl>, carrier <chr>, flight <int>, tailnum <chr>,
## #   origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>, hour <dbl>,
## #   minute <dbl>, time_hour <dtm>
```

Since arr_time is also missing, these are canceled flights.

```
# this is not the tidyverse way:
table(filter(flights, is.na(dep_time))$arr_time, useNA = "always")
```

```
##
## <NA>
## 8255
```

```
# tidyverse way (which we will learn only later in the course)
filter(flights, is.na(dep_time)) %>% count(arr_time)
```

```
## # A tibble: 1 x 2
##   arr_time      n
##   <int> <int>
## 1      NA  8255
```

Exercises 3

1. Sort flights to find the most delayed flights. Find the flights that left earliest

```
# The most delayed flights are found by sorting by dep_delay in descending order.
# There was a flight delayed more than 21 hours.
arrange(flights, desc(dep_delay))
```

```
## # A tibble: 336,776 x 19
##   year month   day dep_time sched_dep_time dep_delay arr_time
##   <int> <int> <int>   <int>         <int>      <dbl>   <int>
## 1  2013     1     9     641             900    1301.    1242
## 2  2013     6    15    1432            1935    1137.    1607
## 3  2013     1    10    1121            1635    1126.    1239
## 4  2013     9    20    1139            1845    1014.    1457
## 5  2013     7    22     845            1600    1005.    1044
## 6  2013     4    10    1100            1900     960.    1342
## 7  2013     3    17    2321             810     911.     135
## 8  2013     6    27     959            1900     899.    1236
## 9  2013     7    22    2257             759     898.     121
## 10 2013    12     5     756            1700     896.    1058
## # ... with 336,766 more rows, and 12 more variables: sched_arr_time <int>,
## #   arr_delay <dbl>, carrier <chr>, flight <int>, tailnum <chr>,
## #   origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>, hour <dbl>,
## #   minute <dbl>, time_hour <dtm>
```

Find the flights that left earliest

```
# If we sort dep_delay in ascending order, we get those that left earliest.
# There was a flight that left 43 minutes early.
arrange(flights, dep_delay)
```

```
## # A tibble: 336,776 x 19
##   year month   day dep_time sched_dep_time dep_delay arr_time
##   <int> <int> <int>   <int>         <int>      <dbl>   <int>
## 1  2013    12     7    2040            2123    -43.     40
## 2  2013     2     3    2022            2055    -33.    2240
## 3  2013    11    10    1408            1440    -32.    1549
## 4  2013     1    11    1900            1930    -30.    2233
## 5  2013     1    29    1703            1730    -27.    1947
## 6  2013     8     9     729             755    -26.    1002
## 7  2013    10    23    1907            1932    -25.    2143
## 8  2013     3    30    2030            2055    -25.    2213
## 9  2013     3     2    1431            1455    -24.    1601
## 10 2013     5     5     934             958    -24.    1225
## # ... with 336,766 more rows, and 12 more variables: sched_arr_time <int>,
## #   arr_delay <dbl>, carrier <chr>, flight <int>, tailnum <chr>,
```

```
## #   origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>, hour <dbl>,
## #   minute <dbl>, time_hour <dtm>
```

2. How could you use `arrange()` to sort all missing values to the start? (Hint: use `is.na()`)

```
arrange(flights, desc(is.na(dep_time)), dep_time)
```

```
## # A tibble: 336,776 x 19
##   year month   day dep_time sched_dep_time dep_delay arr_time
##   <int> <int> <int>   <int>         <int>       <dbl>   <int>
## 1  2013     1     1     NA             1630         NA     NA
## 2  2013     1     1     NA             1935         NA     NA
## 3  2013     1     1     NA             1500         NA     NA
## 4  2013     1     1     NA              600         NA     NA
## 5  2013     1     2     NA             1540         NA     NA
## 6  2013     1     2     NA             1620         NA     NA
## 7  2013     1     2     NA             1355         NA     NA
## 8  2013     1     2     NA             1420         NA     NA
## 9  2013     1     2     NA             1321         NA     NA
## 10 2013     1     2     NA             1545         NA     NA
## # ... with 336,766 more rows, and 12 more variables: sched_arr_time <int>,
## #   arr_delay <dbl>, carrier <chr>, flight <int>, tailnum <chr>,
## #   origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>, hour <dbl>,
## #   minute <dbl>, time_hour <dtm>
```

3. Which flights traveled the longest? Which traveled the shortest?

```
# The longest flights are the Hawaii Air (HA 51) between JFK and HNL (Honolulu) at 4,983 miles.
arrange(flights, desc(distance))
```

```
## # A tibble: 336,776 x 19
##   year month   day dep_time sched_dep_time dep_delay arr_time
##   <int> <int> <int>   <int>         <int>       <dbl>   <int>
## 1  2013     1     1     857             900        -3.    1516
## 2  2013     1     2     909             900         9.    1525
## 3  2013     1     3     914             900        14.    1504
## 4  2013     1     4     900             900         0.    1516
## 5  2013     1     5     858             900        -2.    1519
## 6  2013     1     6    1019             900        79.    1558
## 7  2013     1     7    1042             900       102.    1620
## 8  2013     1     8     901             900         1.    1504
## 9  2013     1     9     641             900      1301.    1242
## 10 2013     1    10     859             900        -1.    1449
## # ... with 336,766 more rows, and 12 more variables: sched_arr_time <int>,
## #   arr_delay <dbl>, carrier <chr>, flight <int>, tailnum <chr>,
## #   origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>, hour <dbl>,
## #   minute <dbl>, time_hour <dtm>
```

We will use later how to use the pipe which is handy for plausibility checking and printing all columns:
We can use the `head()` function to take the first row and `glimpse()` to print all columns:

```
arrange(flights, desc(distance)) %>% head(n=1) %>% glimpse()
```

```
## Observations: 1
## Variables: 19
## $ year           <int> 2013
## $ month          <int> 1
## $ day            <int> 1
```

```
## $ dep_time      <int> 857
## $ sched_dep_time <int> 900
## $ dep_delay     <dbl> -3
## $ arr_time      <int> 1516
## $ sched_arr_time <int> 1530
## $ arr_delay     <dbl> -14
## $ carrier       <chr> "HA"
## $ flight        <int> 51
## $ tailnum       <chr> "N380HA"
## $ origin        <chr> "JFK"
## $ dest          <chr> "HNL"
## $ air_time      <dbl> 659
## $ distance      <dbl> 4983
## $ hour          <dbl> 9
## $ minute        <dbl> 0
## $ time_hour     <dtm> 2013-01-01 09:00:00
```

Alternatively, we can use the `print()` function (`width = Inf` to show all columns)

```
arrange(flights, desc(distance)) %>% print(., n=1, width = Inf)
```

```
## # A tibble: 336,776 x 19
##   year month   day dep_time sched_dep_time dep_delay arr_time
##   <int> <int> <int>   <int>         <int>      <dbl>   <int>
## 1  2013     1     1     857             900        -3.    1516
##   sched_arr_time arr_delay carrier flight tailnum origin dest  air_time
##           <int>      <dbl> <chr>   <int> <chr>   <chr> <chr>   <dbl>
## 1           1530       -14. HA        51 N380HA  JFK   HNL     659.
##   distance hour minute time_hour
##      <dbl> <dbl> <dbl> <dtm>
## 1   4983.    9.    0. 2013-01-01 09:00:00
## # ... with 3.368e+05 more rows
```

Which traveled the shortest?

```
# Apart from an EWR to LGA flight that was canceled, the shortest flights are
# the Envoy Air Flights between EWR and PHL at 80 miles.
arrange(flights, distance)
```

```
## # A tibble: 336,776 x 19
##   year month   day dep_time sched_dep_time dep_delay arr_time
##   <int> <int> <int>   <int>         <int>      <dbl>   <int>
## 1  2013     7    27      NA             106         NA      NA
## 2  2013     1     3    2127             2129        -2.    2222
## 3  2013     1     4    1240             1200        40.    1333
## 4  2013     1     4    1829             1615       134.    1937
## 5  2013     1     4    2128             2129        -1.    2218
## 6  2013     1     5    1155             1200        -5.    1241
## 7  2013     1     6    2125             2129        -4.    2224
## 8  2013     1     7    2124             2129        -5.    2212
## 9  2013     1     8    2127             2130        -3.    2304
## 10 2013     1     9    2126             2129        -3.    2217
## # ... with 336,766 more rows, and 12 more variables: sched_arr_time <int>,
## #   arr_delay <dbl>, carrier <chr>, flight <int>, tailnum <chr>,
## #   origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>, hour <dbl>,
## #   minute <dbl>, time_hour <dtm>
```

Exercises 4

1. Brainstorm as many ways as possible to select `dep_time`, `dep_delay`, `arr_time`, and `arr_delay` from `flights`.

```
select(flights, dep_time, dep_delay, arr_time, arr_delay)
select(flights, starts_with("dep_"), starts_with("arr_"))
select(flights, matches("(dep|arr)_(time|delay)$"))
# using ends_with() doesn't work well since it would return both sched_arr_time and sched_dep_time.
# also the base R subsetting with square brackets works:
flights[, c("dep_time", "dep_delay", "arr_time", "arr_delay")]
```

2. What happens if you include the name of a variable multiple times in a `select()` call?

```
select(flights, year, month, day, year, year)
```

```
## # A tibble: 336,776 x 3
##   year month   day
##   <int> <int> <int>
## 1  2013     1     1
## 2  2013     1     1
## 3  2013     1     1
## 4  2013     1     1
## 5  2013     1     1
## 6  2013     1     1
## 7  2013     1     1
## 8  2013     1     1
## 9  2013     1     1
## 10 2013     1     1
## # ... with 336,766 more rows
```

It ignores the duplicates, and that variable is only included once. No error, warning, or message is emitted.

3. Does the result of running the following code surprise you? How do the select helpers deal with case by default? How can you change that default?

```
# The default behavior for contains() is to ignore case.
select(flights, contains("TIME"))
```

```
## # A tibble: 336,776 x 6
##   dep_time sched_dep_time arr_time sched_arr_time air_time
##   <int>         <int>    <int>         <int>      <dbl>
## 1     517           515      830           819        227.
## 2     533           529      850           830        227.
## 3     542           540      923           850        160.
## 4     544           545     1004          1022        183.
## 5     554           600      812           837        116.
## 6     554           558      740           728        150.
## 7     555           600      913           854        158.
## 8     557           600      709           723         53.
## 9     557           600      838           846        140.
## 10    558           600      753           745        138.
## # ... with 336,766 more rows, and 1 more variable: time_hour <dtm>
```

```
select(flights, contains("TIME", ignore.case = FALSE))
```

```
## # A tibble: 336,776 x 0
```

Exercises 5

1. Currently `dep_time` and `sched_dep_time` are convenient to look at, but hard to compute with because they're not really continuous numbers. Convert them to a more convenient representation of number of minutes since midnight.

Actual departure and arrival times, local time zone. It seems they are not minutes as values between 60 and 99 are missing (but the variable description could be better).

To get the departure times in the number of minutes, (integer) divide `dep_time` by 100 to get the hours since midnight and multiply by 60 and add the remainder of `dep_time` divided by 100.

```
mutate(flights,
  dep_time_mins = dep_time %/% 100 * 60 + dep_time %% 100,
  sched_dep_time_mins = sched_dep_time %/% 100 * 60 + sched_dep_time %% 100) %>%
select(dep_time, dep_time_mins, sched_dep_time, sched_dep_time_mins)
```

```
## # A tibble: 336,776 x 4
##   dep_time dep_time_mins sched_dep_time sched_dep_time_mins
##   <int>      <dbl>         <int>         <dbl>
## 1     517         317.           515           315.
## 2     533         333.           529           329.
## 3     542         342.           540           340.
## 4     544         344.           545           345.
## 5     554         354.           600           360.
## 6     554         354.           558           358.
## 7     555         355.           600           360.
## 8     557         357.           600           360.
## 9     557         357.           600           360.
## 10    558         358.           600           360.
## # ... with 336,766 more rows
```

The clean way is to define a function first and then to reuse it.

```
time2mins <- function(x) {
  x %/% 100 * 60 + x %% 100
}
mutate(flights,
  dep_time_mins = time2mins(dep_time),
  sched_dep_time_mins = time2mins(sched_dep_time)) %>%
select(dep_time, dep_time_mins, sched_dep_time, sched_dep_time_mins)
```

```
## # A tibble: 336,776 x 4
##   dep_time dep_time_mins sched_dep_time sched_dep_time_mins
##   <int>      <dbl>         <int>         <dbl>
## 1     517         317.           515           315.
## 2     533         333.           529           329.
## 3     542         342.           540           340.
## 4     544         344.           545           345.
## 5     554         354.           600           360.
## 6     554         354.           558           358.
## 7     555         355.           600           360.
## 8     557         357.           600           360.
## 9     557         357.           600           360.
## 10    558         358.           600           360.
## # ... with 336,766 more rows
```


2. Find the 10 most delayed flights using a ranking function. How do you want to handle ties? Carefully read the documentation for `min_rank()`.

```
# We want to handle ties by taking the minimum of tied values.
# If three flights are have the same value and are the most delayed,
# we would say they are tied for first, not tied for third or second.
mutate(flights,
       dep_delay_rank = min_rank(-dep_delay)) %>%
filter(dep_delay_rank <= 10) %>%
arrange(dep_delay_rank) %>%
select(dep_delay_rank, everything())

## # A tibble: 10 x 20
##   dep_delay_rank year month   day dep_time sched_dep_time dep_delay
##           <int> <int> <int> <int>   <int>         <int>      <dbl>
## 1             1  2013     1     9     641           900      1301.
## 2             2  2013     6    15    1432          1935      1137.
## 3             3  2013     1    10    1121          1635      1126.
## 4             4  2013     9    20    1139          1845      1014.
## 5             5  2013     7    22     845          1600      1005.
## 6             6  2013     4    10    1100          1900       960.
## 7             7  2013     3    17    2321           810       911.
## 8             8  2013     6    27     959          1900       899.
## 9             9  2013     7    22    2257           759       898.
## 10            10  2013    12     5     756          1700       896.
## # ... with 13 more variables: arr_time <int>, sched_arr_time <int>,
## #   arr_delay <dbl>, carrier <chr>, flight <int>, tailnum <chr>,
## #   origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>, hour <dbl>,
## #   minute <dbl>, time_hour <dtm>
```


Exercises 6

1. Look at the number of canceled flights per day. Is there a pattern? Is the proportion of canceled flights related to the average delay?

```
canceled_delayed <-
  flights %>%
  mutate(canceled = (is.na(arr_delay) | is.na(dep_delay))) %>%
  group_by(year, month, day) %>%
  summarise(prop_canceled = mean(canceled),
            avg_dep_delay = mean(dep_delay, na.rm = TRUE))

ggplot(canceled_delayed, aes(x = avg_dep_delay, prop_canceled)) +
  geom_point() +
  geom_smooth()

## `geom_smooth()` using method = 'loess'
```



```
#> `geom_smooth()` using method = 'loess' and formula 'y ~ x'
```

2. Which carrier has the worst delays? Challenge: can you disentangle the effects of bad airports vs. bad carriers? Why/why not? (Hint: think about flights %>% group_by(carrier, dest) %>% summarise(n()))

```
flights %>%
  group_by(carrier) %>%
  summarise(arr_delay = mean(arr_delay, na.rm = TRUE)) %>%
  arrange(desc(arr_delay))
```

```
## # A tibble: 16 x 2
##   carrier arr_delay
##   <chr>      <dbl>
## 1 F9         21.9
## 2 FL         20.1
## 3 EV         15.8
## 4 YV         15.6
## 5 OO         11.9
## 6 MQ         10.8
## 7 WN          9.65
## 8 B6          9.46
## 9 9E          7.38
## 10 UA         3.56
## 11 US         2.13
## 12 VX         1.76
## 13 DL         1.64
## 14 AA         0.364
```

```
## 15 HA          -6.92
## 16 AS          -9.93
```

```
filter(airlines, carrier == "F9")
```

```
## # A tibble: 1 x 2
##   carrier name
##   <chr>    <chr>
## 1 F9      Frontier Airlines Inc.
```

Frontier Airlines (FL) has the worst delays.

You can get part of the way to disentangling the effects of airports vs. carriers by comparing each flight's delay to the average delay of destination airport. However, you'd really want to compare it to the average delay of the destination airport, after removing other flights from the same airline. But this is beyond the scope of this tutorial.

3. (advanced) For each plane, count the number of flights before the first delay of greater than 1 hour.

```
flights %>%
  arrange(tailnum, year, month, day) %>%
  group_by(tailnum) %>%
  mutate(delay_gt1hr = dep_delay > 60) %>%
  mutate(before_delay = cumsum(delay_gt1hr)) %>%
  filter(before_delay < 1) %>%
  count(sort = TRUE)
```

```
## # A tibble: 3,755 x 2
## # Groups:   tailnum [3,755]
##   tailnum      n
##   <chr>    <int>
## 1 N954UW     206
## 2 N952UW     163
## 3 N957UW     142
## 4 N5FAAA     117
## 5 N38727      99
## 6 N3742C      98
## 7 N5EWAA      98
## 8 N705TW      97
## 9 N765US      97
## 10 N635JB      94
## # ... with 3,745 more rows
```